

# MOHAMMAD ASIF

Data & AI Engineer | MSc Computing (AI) student, Dublin City University  
asifdanish7654@gmail.com | +353 89 274 3006 | Dublin, Ireland | linkedin.com/in/md-asif-dataengineer

## SUMMARY

Data engineer with five years at Cognizant, EY and PwC, mostly building Azure data platforms with Databricks, PySpark, Data Factory and SQL. Over the last two years my work has moved towards GenAI: a RAG assistant for a PwC client, and an agent-based data-quality prototype at EY. I'm now doing an MSc in Computing (AI) at DCU and looking for data or AI engineering roles in Ireland.

## SKILLS

**Data engineering:** Python, PySpark, Spark SQL, SQL, R, Azure Databricks, Azure Data Factory, ADLS Gen2, Delta Lake, medallion architecture, ETL/ELT, data warehousing, RDBMS, HDFS

**GenAI / ML:** Azure OpenAI, LangChain, LangGraph, RAG, vector search, AI agents, MCP, prompt engineering, LLM evaluation, MLflow

**Tools:** Power BI, Azure Key Vault, Git, Agile delivery

## WORK EXPERIENCE

**Data & AI Engineer**, PricewaterhouseCoopers (PwC), Kolkata, India Dec 2025 – Jul 2026

- Owned ingestion for a financial services client: Azure Data Factory pipelines pulling from on-prem SQL Server, Oracle and a few REST APIs into ADLS Gen2, landed as bronze, silver and gold Delta tables. Data retrieval got roughly 40% faster.
- Halved the run time of our slowest Databricks jobs. Most of the gain came from fixing skewed partitions and switching large lookups to broadcast joins.
- Helped build a RAG assistant on Azure OpenAI and Databricks Vector Search that lets the client's teams ask questions about internal policy documents and get answers with page references. My part was the chunking and embedding pipeline, plus evaluating answer quality in MLflow.
- Moved legacy SSIS packages off on-prem servers and onto Data Factory and Databricks.
- Set up MLflow tracking and feature-ready Delta tables for the data science team, working closely with analysts and solution architects on requirements.

**Data & AI Engineer**, Ernst & Young (EY), Kolkata, India Jul 2024 – Dec 2025

- Wrote the PySpark data-quality checks that ran before every release (nulls, duplicates, schema drift, source-to-target reconciliation). Post-release defects dropped by around 40%.
- Extended those checks with a LangGraph agent prototype that profiled Delta tables through an MCP server and drafted new validation rules for an engineer to review.
- Dug into a telecom client's churn data (billing, usage, support tickets) to find out why customers were leaving. The retention changes that followed cut churn by 25%.
- Moved 100+ passwords, keys and connection strings out of pipeline configs and into Azure Key Vault.
- Set up Delta Lake tables with MERGE upserts behind the client's real-time Power BI dashboards.

**Data Engineer**, Cognizant, Bangalore, India Jul 2021 – Jun 2024

- Wrote PySpark and Spark SQL jobs in Databricks to clean up raw CSV, JSON and database extracts for reporting.
- Led pre-release testing, which mostly meant reconciling row counts and values between source and target.
- Ran 20+ internal sessions on Spark and Databricks for new joiners.

## EDUCATION

**MSc in Computing (Artificial Intelligence)**, Dublin City University, Ireland Sep 2026 – 2027

**B.Tech, Information Technology**, Institute of Technical Education and Research, Bhubaneswar, India 2017 – 2021

## CERTIFICATIONS

Microsoft Azure Data Engineer Associate (DP-203) | Azure AI Fundamentals (AI-900) | Azure Fundamentals (AZ-900) | Agentic AI Engineering and The Ultimate Big Data Masters Program (Trendytech)